

Comparing RNA Covariance Models using Zhang Shasha Tree Edit Distance

Bryan Richlinski *University of Kansas* b748r023@ku.edu

Dominic Tassio *University of Kansas* dominic.tassio@ku.edu

Joseph Vinduska *University of Kansas* joseph.vinduska@ku.edu

I. INTRODUCTION

RNA molecules can be grouped into families based on shared characteristics such as structural patterns, sequence similarities, and functional roles. This allows researchers to study them in a more organized and meaningful way. To capture these similarities, a technique known as Multiple Read Alignment (MRA) is often employed, in which several RNA sequences from the same family are aligned together. From this alignment, a model can be constructed that encodes the common features of the family. This model can then be used as a reference to evaluate new RNA reads, providing a measure of how likely it is that a given sequence belongs to the same RNA family.

The chosen model for representing RNA families has often been a Stochastic Markov Model (SMM), which is capable of describing many RNA molecules in terms of both their raw symbolic strings and certain structural configurations that the molecules can take. This model has the same expressive power as a Stochastic Regular Grammar (SRG), meaning it can capture a wide range of sequence and structural patterns. However, it is limited in that it cannot represent more complex structures. These more intricate configurations are referred to as Secondary Structures (SS), and they require more powerful modeling techniques to be fully captured. To address this issue, a Covariance Model (CM) is often used to model these structures. The CM is a form of restricted Stochastic Context Free Grammar (SCFG) which captures a subset of the language of nested strings. These CM's are often represented as tree structures consisting of nodes. The nodes of the tree structure either serve structural purposes, by adding in bifurcation of the structure, or contain states equipped with transition and emission probabilities.

We hypothesize that the structural similarity of these CM trees can be used as a proxy of RNA Family similarity. To this end we have developed an algorithm that utilizes the Zhang Shasha Tree Edit distance algorithm, augmented with node to node comparison that uses a HMM Distance comparison function. This algorithm takes as input two CM's and measures their structural similarity.

We validate our algorithm as finding the smallest edit distance between two Covariance Models in Section II. To validate the efficacy of our approach we compare our results to the clan of RNA families. These clans encode how similar families are to each-other based upon empirical evidence. By comparing the distance between models, to their respective clan we infer a cutoff and validate upon a hold out set. We evaluate this approach in Section V.

II. FORMALISM

A. Zhang-Shasha Algorithm

In our approach to the problem, we decided to use the Zhang-Shasha algorithm for tree edit distance. The Zhang-Shasha algorithm (published by Kaizhong Zhang and Dennis Shasha in 1989) is a dynamic programming algorithm that computes tree edit distance for ordered, labeled trees. These terms are straightforward: in labeled trees, nodes are assigned one of a selection of distinct types (labels), and exchanging two nodes may produce a different tree, if they are not of the same type. Ordered trees, as the name implies, rely on a fixed (ordered) interpretation of their children, usually left to right. Again, swapping a left child and a right child may produce a distinct tree, depending on their own labels and descendants. Fortunately, this fits our scenario well. The nodes in the trees we consider have distinct, meaningful contents, so labels are appropriate. Since the models are RNA-derived, and RNA is read left to right, the notion of order should be applicable as well. There are also practical considerations, as according to Akutsu et. al. (2009), in further work, Zhang and Shasha proved that the general case of tree edit distance for unordered trees is NP-hard.

As a dynamic programming algorithm, the Zhang-Shasha algorithm bears some similarity to other dynamic programming algorithms, such as string edit distance. To calculate tree edit distance, the minimal cost is calculated to transform one tree into another using three basic operations: adding, deleting, or re-labeling a node. However, as trees allow arbitrary connections and are more complicated than strings, additional methods are needed. Two phases are needed to calculate tree edit distance. A dynamic programming table is built up, calculating the cost of changing any sub-tree in the input to any sub-tree in the output, and the results combined for the total tree edit distance, and for each sub-tree, the authors generalize the algorithm to forest edit distance.

By numbering the nodes in each (sub)tree by performing a postorder traversal, children are given lower numbers than their parents, with the leftmost descendant of the root receiving the number 1 and the root being assigned some value n . By selecting a range $[a...b]$, a sub-forest is selected. Using a fixed a and increasing b allows the algorithm to have a base case (a single node) and build up to the entire sub-tree to complete the primary dynamic programming matrix.

Algorithm 1 The Zhang-Shasha Algorithm

Require: Initial tree T_1 and target tree T_2

▷ Setup/Preprocessing

```

1:  $LRK_1 \leftarrow$  all keyroots of  $T_1$ 
2:  $LRK_2 \leftarrow$  all keyroots of  $T_2$ 
3: for each node  $i$  in  $T_1$  and  $T_2$  do
4:    $l[i] =$  leftmost descendant of  $i$ 
5: end for

```

▷ Main Algorithm

```

6: for  $i$  in  $LRK_1$  do
7:   for  $j$  in  $LRK_2$  do
8:      $treedist[i][j] = \text{find\_treedist}(i, j)$ 
9:   end for
10: end for
11: return  $treedist[i][j]$ 

```

One key aspect of this algorithm is that the cost of adding, deleting, or re-labelling nodes is determined by functions that the implementer can provide. As discussed below, this project implements a custom re-labelling cost, determined by the specific HMMs contained in each node. Once the criteria for costs are determined, the algorithm used a simple selection criteria:

$$forestdist[l[i_1]][i][l[j_1]][j] = \min \begin{cases} forestdist[l[i_1]][i-1][l[j_1]][j] + delete_cost(end_1) \\ forestdist[l[i_1]][i][l[j_1]][j-1] + add_cost(end_2) \\ forestdist[l[i_1]][l[i]-1][l[j_1]][l[j]-1] + forestdist[l[i]][i-1][l[j]][j-1] + rename_cost(i, j) \end{cases}$$

i_1 and j_1 are arbitrary ancestors of i and j used to enforce a meaningful range.

To improve the algorithm runtime, the nodes considered can be reduced. Since the final result, tree-to-tree distance, is constructed by subtrees, Zhang and Shasha introduce the notion of Keyroots. Effectively, keyroots are the nodes in the start and end tree where distinct subtrees (and sub-forests) can be selected. Finding them is simple: it is the root node, and any node that has a sibling to its left. It is possible to improve the algorithm by re-using additional work. Since sub-trees and sub-forests are calculated in increasing size, calculating forest edit distance can re-use the tree edit distance for the trees within. This gives updated selection criteria. When operating on a single node, that is, $l[i] = l[i_1]$ and $l[j] = l[j_1]$, we use the selection

$$forestdist[l[i_1]][i][l[j_1]][j] = \min \begin{cases} forestdist[l[i_1]][i-1][l[j_1]][j] + delete_cost(end_1) \\ forestdist[l[i_1]][i][l[j_1]][j-1] + add_cost(end_2) \\ forestdist[l[i_1]][l[i]-1][l[j_1]][l[j]-1] + rename_cost(i, j) \end{cases}$$

And otherwise, re-use the tree distance:

$$forestdist[l[i_1]][i][l[j_1]][j] = \min \begin{cases} forestdist[l[i_1]][i-1][l[j_1]][j] + delete_cost(end_1) \\ forestdist[l[i_1]][i][l[j_1]][j-1] + add_cost(end_2) \\ forestdist[l[i_1]][l[i]-1][l[j_1]][l[j]-1] + treedist(i, j) \end{cases}$$

Note that this algorithm runs in $O(n^4)$ time and $O(n^2)$ space. While the full details are provided in the analysis section, it is worth summarizing quickly. Building the array of all possible tree-to-tree edit distances requires n^2 values, as it is bounded by the number of nodes in the largest tree; and each of these values requires an array of forest-to-forest edit distance, again bounded by the number of nodes in the larger tree. $O(n^2)$ by $O(n^2)$ yields the $O(n^4)$ time complexity, but space remains $O(n^2)$, as only a single instance of the tree-to-tree table and forest-to-forest table are required at once.

B. Hidden Markov Model

A Hidden Markov Model (HMM) consists of a finite set of states X and an alphabet Y . Each state $x \in X$ has an emission probability function, $p(y|x_i)$, that gives the probability of emitting a symbol $y \in Y$ given that the model is in hidden state x_i at time-step i . Each state $x \in X$ has a transition probability function that specifies the probability of moving from the current state to any other state in X at the next observation period, given as $p(x_t|x_{t-1})$.

This model can emit a sequence of observations, $y_1, y_2, \dots, y_T = y_{1:T}$. The likelihood that a given model has produced a particular observation sequence can be computed efficiently using the forward algorithm.

A naive solution to find the probability of producing some sequence $y_{1:T}$ and ending in state x_T can be naively computed by summing the probability of all possible paths which could have emitted the entire sequence and ended in state x . Instead, the forward algorithm uses the Markov property, ie. that the current state only depends on the previous state, to break the problem into optimal subproblems. We can say that $(x_T|y_{1:T})$ only depends upon the solutions to $(x_{T-1}|y_{1:T-1})$ (by the Markov property), and the constant transition / emission probabilities applied to these subproblem solutions.

Algorithm 2 subprocedure: find_treedist

Require: node i in T_1 and node j in tree T_2
Require: LRK_1 , LRK_2 , and $l[x]$ in scope

```

1:  $forestdist[0][0][0][0] \leftarrow 0$ 
     $\triangleright$  Not really a 4-D array: the first and third indices are fixed, so it is actually 2-D
     $\triangleright$  Base case: any sub-forest from  $T_1$  can be reduced to the empty forest

2: for  $end_1$  in  $range(l[i], i)$  : do //inclusive range
3:    $forestdist[l[i]][end_1][0][0] = forestdist[l[i]][end_1 - 1][0][0] + delete\_cost(i)$ 
4: end for
     $\triangleright$  Base case: any sub-forest in  $T_2$  can be constructed from the empty forest

5: for  $end_2$  in  $range(l[j], j)$  : do //inclusive range
6:    $forestdist[0][0][l[j]][end_2] = forestdist[0][0][l[j]][end_2 - 1] + insert\_cost(j)$ 
7: end for

8: for  $end_1$  in  $range(l[i], i)$  : do
9:   for  $end_2$  in  $range(l[j], j)$  : do
10:    if  $l[end_1] == l[i]$  and  $l[end_2] == l[j]$  then  $\triangleright$  forestdist completed
11:       $result = \min \begin{cases} forestdist[l[i]][end_1 - 1][l[j]][end_2] + delete\_cost(end_1) \\ forestdist[l[i]][end_1][l[j]][end_2 - 1] + add\_cost(end_2) \\ forestdist[l[i]][end_1 - 1][l[j]][end_2 - 1] + rename\_cost(end_1, end_2) \end{cases}$ 
12:       $treedist[end_1][end_2] = result$ 
13:      return  $result$ 
14:    else
15:       $result = \min \begin{cases} forestdist[l[i]][end_1 - 1][l[j]][end_2] + delete\_cost(end_1) \\ forestdist[l[i]][end_1][l[j]][end_2 - 1] + add\_cost(end_2) \\ forestdist[l[i]][l[end_1] - 1][l[j]][l[end_2] - 1] + treedist[end_1][end_2] \end{cases}$ 
16:       $forestdist[l[i]][end_1][l[j]][end_2] = result$ 
17:    end if
18:  end for
19: end for

```

Algorithm 3 Forward Algorithm for HMM

```

1: Initialize
2:    $t = 0$ 
3:   transition probabilities:  $p(x_t | x_{t-1})$ 
4:   emission probabilities:  $p(y_t | x_t)$ 
5:   observations:  $y_{1:T}$ 
6:   prior probability:  $\alpha(x_0)$ 

7: For  $t = 1$  to  $T$ :
8:    $\alpha(x_t) = p(y_t, x_t) \sum_{x_{t-1}} p(x_t | x_{t-1}) \alpha(x_{t-1})$ 

9: Return:  $p(y_{1:T}) = \sum_{x_T} \alpha(x_T)$ 

```

In algorithm 3, we present the forward algorithm. After initializing the model parameters and prior probabilities, the algorithm iterates through each time step $t = 1$ to $t = T$. At each iteration the algorithm updates the probability of being in a given state and emitting the current symbol in the sequence. It does this by summing over all probabilities from the previous time step, each weighted by their associated transition probabilities, and multiplying the sum by the emission probability of the current observation given the state currently being calculated for.

The final result of this step is the probability of being in state x at time T , given the entire sequence of observations up to that point. This calculation is repeated for each hidden state at each time in the sequence. To calculate the total probability of the HMM emitting the sequence, regardless of final state, the algorithm will sum over all final state probabilities.

C. Distance Calculation

To compare two Hidden Markov Models H_1 and H_2 , our comparison algorithm generates several sequences of observations from each model. Let these sequences be labeled $h_{11}, h_{12}, \dots, h_{1n}$. These sequences are then used to approximate the Jensen Shannon (JS) divergence between H_1 and H_2

$$D(H_1, H_2) \approx \frac{1}{2n} \left(\sum_{i=1}^n p(h_i | H_2) * \log_2 \frac{p(h_i | H_2)}{M} + \sum_{i=1}^n p(h_i | H_1) * \log_2 \frac{p(h_i | H_1)}{M} \right). \quad (1)$$

In equation 1 $p(h | H)$ denotes the probability of observing sequence h under model H , computed using the forward algorithm. M is a mixture model of the probabilities, here given by $0.5 * (p(h | H_1) + p(h | H_2))$. This is merely an approximation of the JS Divergence, as a true measure would require full knowledge of the probability distributions of H_1 and H_2 , but this provides an approximation.

III. ANALYSIS

A. ZSS Complexity

The Zhang-Shasha algorithm runs with time complexity $O(n^4)$ and space complexity $O(n^2)$, where n is bounded by the number of nodes in the larger tree. More specifically, Zhang and Shasha state that the algorithm requires $O(|T_1| * |T_2| * \min(\text{depth}(T_1), \text{leaves}(T_1)) * \min(\text{depth}(T_2), \text{leaves}(T_2)))$ in time and $O(|T_1| * |T_2|)$ in space. (Zhang and Shasha, 1989) As described in the formalism section, this is due to the two-stage dynamic programming approach. Tree edit distances are computed for each subtree in T_1 to each subtree in T_2 (both clearly bounded by n , the number of nodes in the larger tree) which requires $O(n^2)$ time and space. However, for each of these calculations, further dynamic programming is required. Each sub-tree edit distance is computed via sub-forest edit distance, again, from any possible sub-forest from the T_1 subtree to any subforest in the T_2 subtree. Since the subtrees are bounded by the n , and the sub-forests bounded by the size of the subtrees, this inner calculation is also $O(n^2)$ for every one of the tree-to-tree cells which requires it. Combining the $O(n^2)$ outer calls with the $O(n^2)$ inner calls produces the $O(n^4)$ time complexity. Since only one table is needed for the outer calls, and only one inner call is active at once, the space complexity increases linearly, and $O(n^2) + O(n^2)$ remains $O(n^2)$.

More modern versions of this algorithm exist, with slightly better ($O(n^3)$) time complexity. However, these algorithms are much more difficult to implement, and Paaßen (2019) cites Pawlik and Augsten (2011) in explaining that "... the practical runtime complexity of the original TED algorithm is still competitive for balanced trees, such that we regard it as a good choice in many practical scenarios." Further, as the number of bifurcation nodes in a CM is relatively few, the number of keyroots found by the algorithm will also be relatively small, further reducing the work needed in our specific scenario.

B. HMM Complexity

The forward algorithm has time complexity $O(nm^2)$, where n is the length of the observation sequence and m is the number of hidden states. For a model with m states, at time step t , the algorithm computes the product of the prior probability of being in a given state at time $t-1$ with the probability of transitioning to the current state under consideration, this is done for all m prior probabilities and all m states considered in the current time step. This yields the m^2 term. As this calculation must be done for all characters in the emission sequence we arrive at the n term.

While the forward algorithm is a dynamic algorithm, due to the nature of continuous probabilities it is very unlikely that subproblems will have repeated solutions, therefore no memoization table is used. Consequently, the space complexity of the forward algorithm is $O(m)$, m being the number of hidden states in the model.

When incorporated into the ZSS algorithm the forward algorithm and JS distance runs only once for each pair of nodes being compared, due to memoization. This means that when comparing two trees the distance measures are run a maximum of $N_1 * N_2$ times, where N_i is the number of nodes in the trees resp. This means that the distance measure contributes a fixed time complexity of $O(N^2 * nm^2)$. We can simplify the n term to be the largest string seen, making that a constant, and we can simplify the m term to be the largest HMM seen. This simplifies this to an addition of a $O(N^2 * m^2)$.

The label distance algorithm requires we use the HMM to generate multiple random, but statistically likely, sequences, each of indeterminate length. As this is a random process, a worst-case time complexity cannot be specified for this generation step without fixing the length of the generated strings. However, if the length of a string is known, the complexity of its generation is linear in that length. Therefore, the cost of string generation can be folded into the overall complexity of the forward algorithm. Being unable to specify the length of the sequences in advance does mean that empirical measurements of time complexity may deviate from the time complexity presented in the ZSS algorithm, due to the confounding n term.

IV. IMPLEMENTATION

The CM distance algorithm is implemented in Python. We make use of the `hmmlearn` library to store Hidden Markov Models, to emit sequences of observations, and to apply the forward algorithm for likelihood calculations. For the ZSS algorithm, we employ the `zss` package to perform the Zhang-Shasha tree edit distance, augmented with our custom distance metric to tailor the comparison to our models.

A. Node HMM Model Augmentations

In contrast to the formal descriptions provided in Formalisms, the concrete implementation requires structural adjustments to the Hidden Markov Models. Specifically, each model is augmented with a single entry state and a final exit state. The reason for the initial incoming transition is that, while we represent each node as an HMM, this is a simplification. A node may have several incoming and several outgoing transitions, potentially in a non-functional, many-to-many relationship. To account for this, these incoming and outgoing transitions are collected and normalized into one incoming and one outgoing state resp. This ensures that every sequence begins from a unique entry point and terminates in a unique exit point.

As well, in the concrete models there are hidden states that are silent. The forward algorithm cannot directly handle the potentially infinite paths that arise from such ϵ -transitions, so these states must be removed. Eliminating silent states while preserving the transition and emission probabilities of the non-silent states requires an $O(n^3)$ preprocessing step, where n is the number of states in the model. This transformation is performed once for each node when the model is parsed from file.

Alternatively, we could have chosen to force every silent state to emit a self-identifying character. This modification may shift the comparison between HMMs toward a more structural rather than emission-based distance measure. Further research could evaluate the effect of this approach.

B. Distance Metric Augmentations

Because the ZSS algorithm must also account for insertions and deletions, we require a rule for comparing a node to its absence. We handle this by comparing the observed sequences emitted from the node, to the probability that it emits the empty string.

For this we calculate the Jensen Shannon Divergence of the probabilities of the node's observed emissions compared to the probabilities that it emitted the empty string. As opposed to the probability of a different model emitting said observation.

Additionally, certain nodes in the concrete model cannot be represented as Hidden Markov Models. Specifically, nodes labeled BIF and END act as structural markers rather than probabilistic models. Any time these nodes are compared to like named nodes they will return a similarity of 0.0 (perfect match), and return the value 1 when compared to non-similar nodes. $\log_2(2) = 1$ is the maximum value of the JSD function when using base two for the log calculations.

C. Memoization

To reduce redundant computation, many values are saved during parsing and model construction. All sample strings required for label comparison are generated once and stored, together with the probability that the node produced it, given by the forward algorithm.

Furthermore, when comparing two CM's using the ZSS algorithm, the distance between any two nodes is saved after their first computation. Subsequent distance measures between these two nodes are retrieved directly from this cache, eliminating repeated evaluations and improving overall efficiency.

D. Hyperparameters

The only hyperparameter in our distance function is the number of strings generated for the Node distance comparison. currently this is set at 3, larger numbers may yeild more accurate results.

V. EVALUATION

Having implemented this approach for comparing covariance models, we need a way to evaluate its performance. We explore two methods to evaluate how well our approach compares covariance models. First, is an extremely naïve statistical analysis that will use the RNA family's identified clan to compare within- and between-group similarity. Second, is a more nuanced approach that utilizes concrete sequence comparison as a base truth for comparing the models.

Due to time constraints, we only fully implemented the first comparison method. We will still describe the second method as it is useful for context and potential future work.

For this first method, we utilized the clan classification given to some groups of RNA families. Naïvely, we assume that the distance of covariance models within a clan is on average smaller than the distance of covariance models between clans. However, using clans to perform within- and between-group analysis is flawed for several reasons. Primarily, the clans are not strictly organized based on structure. As our approach implements a structural comparison between covariance models, comparing the distance of models within clans and models between clans will not necessarily be representative. Another flaw being the variability in the number of identified RNA families in a clan. As we have an uneven distribution of potential within- and between-clan comparisons, our comparison will need to be resilient to that bias in the dataset. Despite these flaws, we still attempt this comparison method.

Permutational multivariate analysis of variance (PERMANOVA) is a good statistical test for this comparison. Compared to ANOVA, which compares the means of groups, PERMANOVA uses a distance measurement between values. This matches perfectly with our need to evaluate our distance comparison and mitigates the impact of our disparate group sizes.

We arbitrarily selected 40 covariance models from 10 clans for use in our evaluation. This gives 780 combinations of families to compare. On an M4 MacBook Pro with 24GB of RAM, this took approximately 3 minutes to complete. Given the distance from these combinations, we construct the distance matrix necessary for PERMANOVA. Running this statistical test resulted in a test statistic of approximately **5.1** and a p-value of approximately **0.002**. The test statistic describes the similarity of the within- and between-group comparisons. We will discuss the implications of these results in Section VI.

Our second method of evaluation is to use the comparison of RNA sequences belonging to an RNA family as a baseline truth to measure our comparison of covariance models. As the comparison of RNA sequences is a structural comparison, the flaw of using clans would not be an issue. If our method of comparing covariance models is an effective measurement, we would expect there to be a correlation for the distance measured between RNA sequences of different families and distance between covariance models. Unfortunately, due to time constraints and our inexperience in bioinformatics, we were unable to perform this analysis.

VI. DISCUSSION

From our evaluation of our comparison approach using PERMANOVA, we have found that using the Zhang-Shasha tree edit distance algorithm and our HMM distance cost function produced results that were statistically significant, but not necessarily meaningful. The test statistic of approximately **5.1** means that there is significant difference in the within- and between-group distances. Further, as our p-value is **0.002**, our results are likely statistically significant.

There are threats to our approach. While Zhang-Shasha is a competitive algorithm for tree edit distance of ordered and labeled trees, it is possible that our chosen ordering of nodes from the covariance model negatively impacted our results. It is also possible that the cost function we implemented to determine the cost of deleting, inserting, and re-labeling nodes was not optimal. Our HMM fragment to HMM fragment is a probabilistic process involving produced sequence likely hood and could be potentially tuned for better or worse results. Without a rigorous method of evaluation, these individual aspects of our approach are difficult to properly evaluate and potentially tune.

Another potential threat is our small dataset. We only considered 40 covariance models and 10 clans. At the time, there are 4,170 RNA families and 147 clans on Rfam, meaning that there is a much larger dataset we could reach for.

Using all 4,170 RNA families (assuming each family on Rfam has a covariance model) would mean just under 8.7 million comparisons to make. It is unlikely we could have completed our approach on the entire dataset in a reasonable time frame using the consumer hardware available to us. Despite this limitation, it is possible that the PERMANOVA test on this larger dataset could produce different results. However, due to the previously discussed issue with using clans as a category, this approach would likely still be flawed.

The second method of using correlation of RNA sequences comparison to covariance model comparison is a much more promising avenue for fairly evaluating our approach. RNA sequence comparison or alignment is a well studied problem in bioinformatics, with a number of approaches and techniques. Given enough time, we likely could have repurposed an existing tool for this. We regret that we were unable to conduct this evaluation in time to include in this paper.

One benefit of our approach is the actual run time for comparing two models. While we have no concrete data for any one particular run, being able to perform 780 comparisons in approximately 3 minutes on consumer hardware represents a reasonable run time. Due to our implementation being written in Python and utilizing no amount of parallelization or caching between comparisons, it is likely that there are significant performance gains to be made. As each comparison is independent of any other comparison, parallelization would be an easy optimization that would likely see time improvements, even on consumer hardware.

Overall, we find our approach to have significant potential despite the potential flaws in our evaluation. Further work is needed to verify our approach and confirm the promising results we have.

Originally we presented a negative result to this approach. This was due a bug in the HMM comparison.

VII. RELATED WORK

VIII. CONCLUSION

We present our approach for structural comparison of RNA family covariance models using tree edit distance and HMM to HMM comparison. The Zhang-Shasha algorithm is a competitive and robust method of computing the distance between two ordered and labeled trees in a reasonable time complexity for the task. Our HMM to HMM comparison leverages the internal structure of nodes within a covariance model to produce a cost function that can be used in the tree edit distance algorithm. Together, we find these two approaches to be an interesting, if not promising, way to compare covariance models.

Despite the flaws in our evaluation of this approach, we believe that utilizing a different method of evaluation would allow us to better scrutinize this approach and present ways to improve the details of our HMM to HMM comparison. It is unfortunate that we were unable to perform a better evaluation before the writing of this paper. This leaves significant future work open for developing and evaluating this approach.

Ultimately, we find that this approach was a worthwhile first attempt at using tree edit distance for comparison of covariance models.

IX. DISCLAIMER

This report was written with assistance from an llm working in a strictly editorial capacity. All content and ideas are the original work of the authors.